

Applied Machine Learning

Lecture 5-1: Optimization in machine learning

Selpi (selpi@chalmers.se)

The slides are further development of Richard Johansson's slides

February 7, 2020

Overview

Optimization in machine learning

SGD for training a linear regression model

introduction to PyTorch

regularized linear regression models


Objective functions for machine learning

- ▶ in this lecture, we'll discuss **objective functions** for our models
- ▶ training a model = minimizing (or maximizing) this objective
- ▶ often, the objective consists of a combination of
 - ▶ a **loss function** that measures how well the model fits the training data
 - ▶ a **regularizer** that measures how simple/complex the model is

minimizing squared errors

- ▶ in the **least squares** approach to linear regression, we want to minimize the mean (or sum) of **squared errors**


$$f(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2$$



- ▶ From calculus: if a “nice” function has a maximum or minimum, then the derivative will be zero there

the gradient

- ▶ the multidimensional equivalent of the derivative is called the **gradient**
- ▶ if f is a function of n variables, then the gradient is an n -dimensional vector, often written $\nabla f(x)$
- ▶ intuition: the gradient points in the uphill direction



- ▶ again: the gradient is zero if we have an optimum
- ▶ this intuition leads to a simple idea for finding the minimum:
 - ▶ take a small step in the direction opposite to the gradient
 - ▶ repeat until the gradient is close enough to zero

Gradient descent, pseudocode

- ▶ the same thing again, in pseudocode:
 1. set x to some initial value, and select a suitable step size η
 2. compute the gradient $\nabla f(x)$
 3. if $\nabla f(x)$ is small enough, we are done
 4. otherwise, subtract $\eta \cdot \nabla f(x)$ from x and go back to step 2
- ▶ conversely, to find the maximum we can do **gradient ascent**: then we instead add $\eta \cdot \nabla f(x)$ to x

Overview

Optimization in machine learning

SGD for training a linear regression model

introduction to PyTorch

regularized linear regression models

stochastic gradient descent (SGD)

- ▶ in machine learning, our objective functions are often defined as sums over the whole training set, such as the least squares loss function

$$f(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2$$

- ▶ if N is large, it would take time to use gradient descent to minimize the loss function
 - ▶ because gradient descent takes into account all the instances in the training set
- ▶ **stochastic gradient descent**: simplify the computation by computing the gradient using just a small part
 - ▶ in the extreme case, a **single** training instance
 - ▶ or **minibatch**: a few instances

SGD: pseudocode

1. set w to some initial value, and select a suitable step size η
2. select a single training instance x
3. compute the gradient $\nabla f(w)$ **using x only**
4. if we are “done”, stop
5. otherwise, subtract $\eta \cdot \nabla f(w)$ from w and go back to step 2

SGD: pseudocode

1. set w to some initial value, and select a suitable step size η
2. select a single training instance x
3. compute the gradient $\nabla f(w)$ **using x only**
4. if we are “done”, stop
5. otherwise, subtract $\eta \cdot \nabla f(w)$ from w and go back to step 2

$w = (0, \dots, 0)$

for $(x_i, y_i) \dots$

 compute gradient ∇f_i for current instance (x_i, y_i)

$w = w - \eta \cdot \nabla f_i(w)$

return w

when to terminate SGD?

- ▶ simple solution: fixed number of iterations
- ▶ or stop when we've seen no improvement for some time
- ▶ or evaluate on a held-out set: **early stopping**

applying SGD to the least squares loss

- ▶ the least squares loss function:

$$f(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2$$

- ▶ let's consider just a single instance:

$$f_i(\mathbf{w}) = (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2$$

- ▶ for this instance, the gradient of the least squares loss with respect to \mathbf{w} is

$$\nabla f_i(\mathbf{w}) = 2 \cdot (\mathbf{w} \cdot \mathbf{x}_i - y_i) \cdot \mathbf{x}_i$$

plugging it back into SGD \Rightarrow Widrow-Hoff again!

```
w = (0, ..., 0)
for (xi, yi) ...
     $\nabla f_i(\mathbf{w}) = 2 \cdot \text{error} \cdot \mathbf{x}_i$ 
     $\mathbf{w} = \mathbf{w} - \eta \cdot \nabla f_i(\mathbf{w})$ 
return w
```

Overview

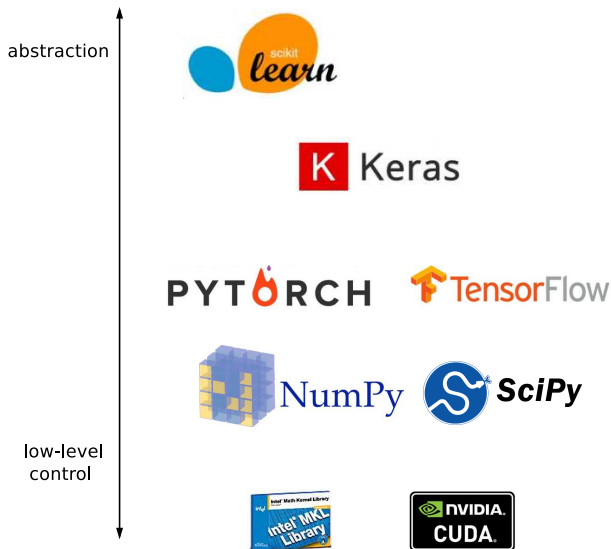
Optimization in machine learning

SGD for training a linear regression model

introduction to PyTorch


regularized linear regression models

the Python machine learning ecosystem (selection)



finding and installing PyTorch

- ▶ easy to install using conda
- ▶ <https://pytorch.org/>



KEY FEATURES & CAPABILITIES

[See all Features >](#)

Hybrid Front-End

A new hybrid front-end seamlessly transitions between eager mode and graph mode to provide both flexibility and speed.

Distributed Training

Scalable distributed training and performance optimization in research and production is enabled by the torch.distributed backend.

Python-First


Deep integration into Python allows popular libraries and packages to be used for easily writing neural network layers in Python.

Tools & Libraries

A rich ecosystem of tools and libraries extends PyTorch and supports development in computer vision, NLP and more.

what is the purpose of PyTorch?

- ▶ on a high level: a library for prototyping ML models
- ▶ on a low level: NumPy-like with **automatic gradients**
- ▶ easy to move computations to a GPU (if available)



[[source](#)]

Overview

Optimization in machine learning


SGD for training a linear regression model

introduction to PyTorch

regularized linear regression models

recall: the fundamental tradeoff in machine learning


- ▶ **goodness of fit**: the learned model should describe the examples in the training data
- ▶ **regularization**: the model should be simple



[Image source]

- ▶ the least squares loss just takes care of the first part!

$$f(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2$$



what does it mean to “keep the model simple”?

- ▶ concretely, how can we add regularization to the linear regression model?
- ▶ the most common approach is to add a term that **keeps the weights small**
- ▶ for instance, by penalizing the squared length (norm) of the weight vector should be small:

$$\|\mathbf{w}\|^2 = w_1 \cdot w_1 + \dots + w_n \cdot w_n = \mathbf{w} \cdot \mathbf{w}$$

- ▶ this is called a **L_2 regularizer** (or a L_2 penalty)

combining the pieces

- ▶ we combine the loss and the regularizer:

$$\begin{aligned} \frac{1}{N} \sum_{i=1}^N \text{Loss}(\mathbf{w}, \mathbf{x}_i, y_i) + \alpha \cdot \text{Regularizer}(\mathbf{w}) \\ = \frac{1}{N} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2 + \alpha \cdot \|\mathbf{w}\|^2 \end{aligned}$$

- ▶ in this formula, α is a “tweaking” parameter that controls the tradeoff between loss and regularization

ridge regression

- ▶ the combination of least squares loss with L_2 regularization is called **ridge regression**

$$\frac{1}{N} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2 + \alpha \cdot \|\mathbf{w}\|^2$$



- ▶ in scikit-learn: `sklearn.linear_model.Ridge`
- ▶ (unregularized least squares: `sklearn.linear_model.LinearRegression`)

discussion

$$\frac{1}{N} \sum_{i=1}^N (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2 + \alpha \cdot \|\mathbf{w}\|^2$$

- ▶ what happens if we set α to 0?
- ▶ what happens if we set α to a huge number?

why “ridge”?



[[Image source](#)]

SGD for the ridge regression model

$$\frac{1}{N} \sum (\mathbf{w} \cdot \mathbf{x}_i - y_i)^2 + \alpha \cdot \|\mathbf{w}\|^2$$

- ▶ the gradient with respect to a single training instance is

$$\nabla f_i(\mathbf{w}) = 2 \cdot \text{error} \cdot \mathbf{x}_i + 2 \cdot \alpha \cdot \mathbf{w}$$

$\mathbf{w} = (0, \dots, 0)$

for (\mathbf{x}_i, y_i) in the training set

$g = \mathbf{w} \cdot \mathbf{x}_i$

error = $g - y_i$

$\mathbf{w} = \mathbf{w} - \eta \cdot (\text{error} \cdot \mathbf{x}_i + \alpha \cdot \mathbf{w})$

return \mathbf{w}

the Lasso and ElasticNet models

- ▶ another common regularizer is the L_1 norm

$$\|\mathbf{w}\|_1 = |w_1| + \dots + |w_n|$$

- ▶ the L_1 norm gives the **Lasso** model
- ▶ combination of L_1 and L_2 gives the **ElasticNet** model
- ▶ in scikit-learn:
 - ▶ `sklearn.linear_model.Lasso`
 - ▶ `sklearn.linear_model.ElasticNet`

why does Lasso give zero coefficients?




FIGURE 3.11. Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions $|\beta_1| + |\beta_2| \leq t$ and $\beta_1^2 + \beta_2^2 \leq t^2$, respectively, while the red ellipses are the contours of the least squares error function.

► figure from [Hastie et al. book](#)